Page 1: Introduction to RAG
Retrieval-Augmented Generation combines a retriever with an LLM.
It lets the model answer using your own documents.

Page 2: Chunking
Documents are split into smaller pieces called chunks.
Chunking strategies include page-wise, character-wise,
recursive and semantic chunking.

Page 3: Retrieval
At query time, the most relevant chunks are retrieved and passed to the LLM as context.